



Department of Information Systems  
and Computer Science

# Data Representation

CS 114





## Text

- How is text represented
- Probabilistic Text
- Sequential Text

## Images

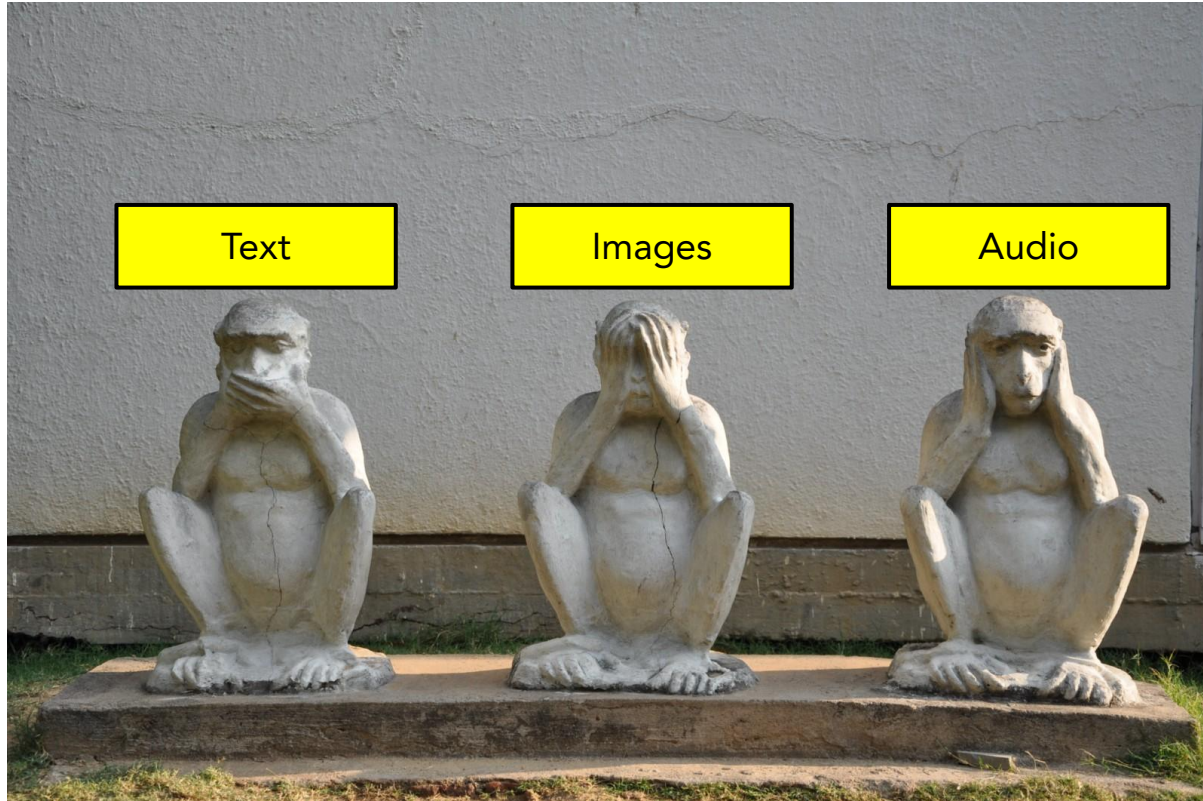
- How are images represented
- Pixel Values
- Convolutions
- Dealing with Video

## Audio

- How is audio represented
- Signal Processing



# Modes of Data





Text



- Text data can be seen as a sequence of alphanumeric characters in its most granular form
- Treating it as a feature vector poses a challenge because of granularity
  - Letter per letter sequence is too noisy
  - Words are too general without context
  - Sentences are hard to comprehend without granularity
- Text data is often characterized as sparse data
- We often first have to define a vocabulary (letter-wise / word-wise / sentence-wise)



# Text and Probability

- An inference to a probabilistic outcome  $P(y | x)$  is not scalable when dealing with sparse data
- Bayesian Inference:

$$P(A|B) = (P(B|A) P(A)) / P(B)$$

$$P(y|x) = (P(x|y) P(y)) / P(x)$$

$$P(\text{"b will occur"} | \text{"a"}) = P(\text{"a occurrences"} | \text{"b"}) P(\text{"b instances"}) / P(\text{"a"})$$

$$\begin{aligned} P(\text{"b will occur"} | \text{"a"}) &= P(\text{"a occurrences"} | \text{"b"}) P(\text{"b instances"}) \\ P(\text{"c will occur"} | \text{"a"}) &= P(\text{"a occurrences"} | \text{"c"}) P(\text{"c instances"}) \end{aligned}$$



# Text and Probability

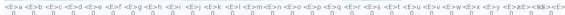
- Bayesian inference on sparse data deals with sparse multiplication operations due to Joint Probability
- Computation on bayesian inference even on just the numerator will be performed over and over again for every instance of "y"

```
P("b will occur" | "a") = P("a occurrences" | "b") P("b instances")  
P("c will occur" | "a") = P("a occurrences" | "c") P("c instances")
```

This is only for the possibilities of "b" vs "c".  
What about for a vocabulary of 10,000 possibilities?

# Text Data for Name Generation

- Bigram Model: Learned weights  $W$  is a count probability of two consecutive letters. By simply increasing the window size of previous letters, we run into computationally expensive probabilities.



Count Occurrences: { "a  $\rightarrow$  b": 10, "a  $\rightarrow$  c": 5, "a  $\rightarrow$  d": 6 }

$\{ "a \rightarrow b": 10, "a \rightarrow c": 5, "a \rightarrow d": 6 \}$

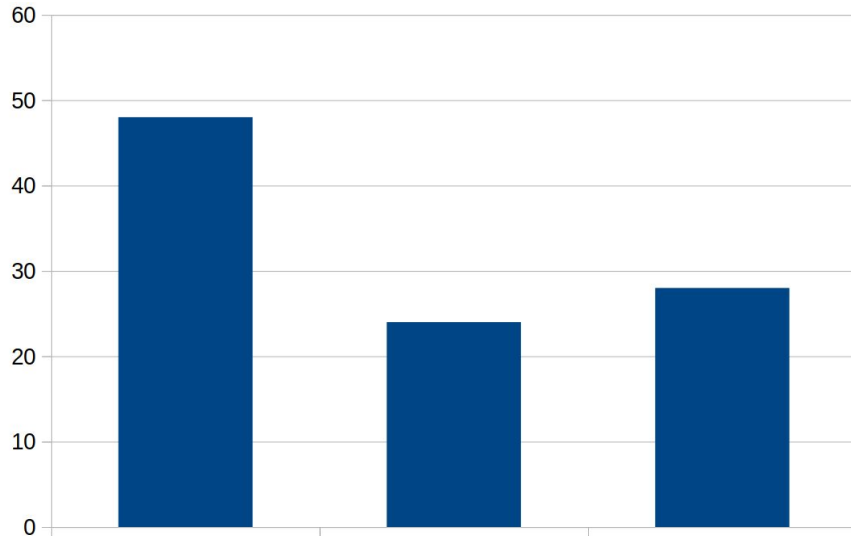
{ “a  $\rightarrow$  b”: 48%, “a  $\rightarrow$  c”: 24%, “a  $\rightarrow$  d”: 28% }





# Problem with Probability

- Text Generation: Randomly sample an item from the probability distribution with an assumption that expresses a multinomial distribution



There is a 48% chance that a “b” comes after an “a”, a 24% chance that a “c” comes after an “a” and a 28% chance that a “d” comes after an “a”. When randomly sampled, these weighted parameters are applied to generate the next letter.

$$\mathbf{f}(\mathbf{x}) = \mathbf{x} * \mathbf{W}$$

{ “a → b”: 48%, “a → c”: 24%, “a → d”: 28% }



# Usefulness of Probability

- Modeling the outcome of the next word or token
- Determine if a model  $f(x)$  can generate the next letter given a previous (or multiple previous letters):

**Vocabulary:** ["a", "b", "c"]

A model is sure to get an "a" next if the probability outcome is:

[1, 0, 0]

*\*Notice that all values sum up to 1 or 100%*

**$f(\text{"b"}) = x * W$**

For some optimized  $W$ , when multiplied to  $x$  we get some outcome. **However, we are not guaranteed that  $W$  will yield a probability distribution (that is, all values sum to 1).**

Example:

[0.52, 0.57, 0.58]



- A mathematical function that takes a vector of arbitrary values and returns a probability distribution over it.

$$s(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

[0.52, 0.57, 0.58]

**$s(0.52) = 1.682 / (1.682 + 1.768 + 1.786) = 0.32$**

**$s(0.57) = 1.768 / (1.682 + 1.768 + 1.786) = 0.34$**

**$s(0.57) = 1.768 / (1.682 + 1.768 + 1.786) = 0.34$**



# Embeddings

- A way to represent words or phrases or sequences as a vector of numbers
- Instead of treating each “token” in a language as possible  $x$  input, we could consider it compressed into an embedding that could also represent other tokens of similar nature
- Objective: Turn a sparse (raw) input into something more compact and understandable. Figure out some  $W$  that generalizes the representation.
- Therefore, an embedding is just a smaller representation of some input transformed by  $W$ .



# Embeddings

- Why are embeddings relevant in text data? Let's take a look at an example:

The dog ate my homework

The cat ate some catfood

**What word can we substitute to make this sentence make sense?**

The \_\_\_\_\_ crossed the road



# Embeddings

The dog ate my homework

Some value to say that “The” refers to “dog” which is an animal

Some animal usually comes before a verb like “ate”

“my homework” refers to some subject

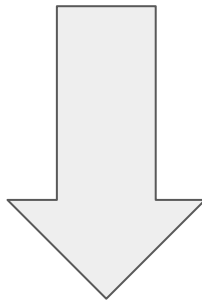
Embedding: Some set of values that strongly relates to:

[refer, animal, verb, subject]



# Embeddings

The cat ate some catfood



**Some magical structure  $W$  to  
create the embeddings**

```
f(["The", "cat", "ate", "some", "catfood"]) = x * W  
  
= [refer, animal, verb, subject]  
= [P(refer), P(animal), P(verb), P(subject)]
```



# Embeddings

- Given some learned  $W$  to represent embeddings and outcomes (and everything in between for that matter), we can then generate the embedding that most likely will complete the following:

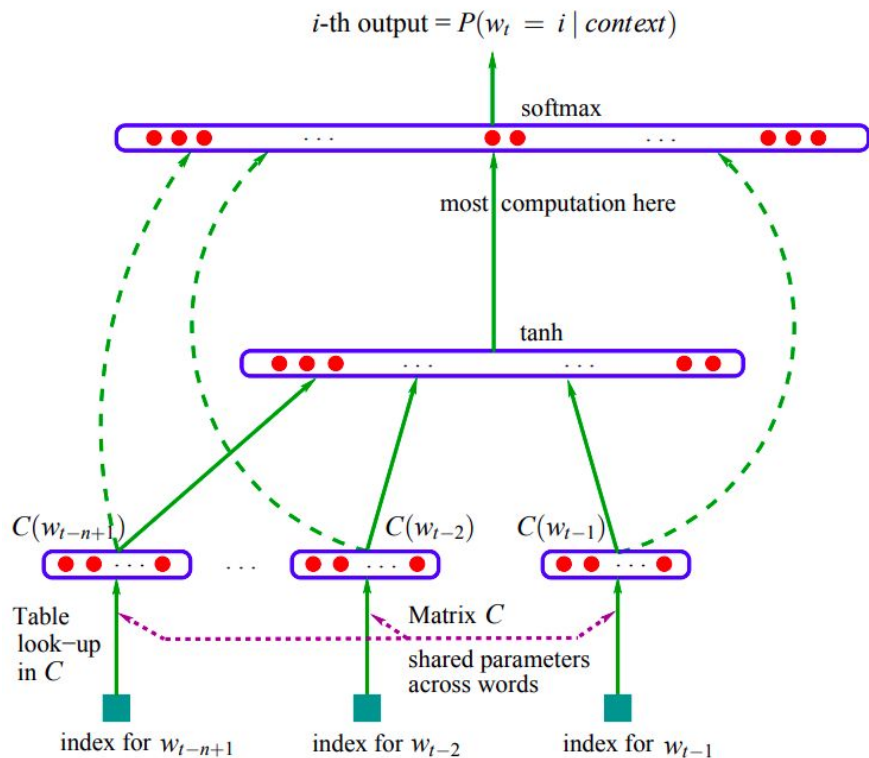
The \_\_\_\_\_ crossed the road

**Magical structure  $W$  should then most likely associate the missing piece with an animal given that most of the parts of the sentence would express similar embedding values. We could sample it from our vocabulary if we know what  $W$  looks like.**





# Bengio's MLP



**Softmax layer representing the next part of the sequence that is token at  $t + 1$**

**Hidden layer that activates  $[-1, 1]$  with the tanh function.**

**An embedding of size 30 that represents each of the input token with shared  $W$  in between**

**A set of 3 tokens to represent the input.  
Sparsely, input size is 17k.**



# Comparing Text

- How good is my generated word or sentence?
- COSINE Similarity: If text A and text B can be turned into some form of embeddings, we can apply a difference function to see how close they are to each other



# Example

I like programming and pizza.

I like pineapple in pizza.

How similar are these  
sentences?



# Vocabulary

- I
- like
- programming
- pineapple
- and
- in
- pizza



# Word Count

Vocabulary	Sentence 1	Sentence 2
I	1	1
like	1	1
programming	1	0
pineapple	0	1
and	1	0
in	1	1
pizza	1	1



# Cosine Similarity

- Similarity measure between two non-zero vectors
- Takes the inner product space that measures the cosine of the angle between them

$$\cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$



- Sentence 1 = [1, 1, 1, 0, 1, 1, 1]
- Sentence 2 = [1, 1, 0, 1, 0, 1, 1]

$$A \text{ dot } B = (1*1) + (1*1) + (1*0) + (0*1) + (1*1) + (1*1) = 4$$

$$||A|| = \text{sqrt}(1^2 + 1^2 + 1^2 + 0^2 + 1^2 + 1^2 + 1^2) = 2.4495$$

$$||B|| = \text{sqrt}(1^2 + 1^2 + 0^2 + 1^2 + 0^2 + 1^2 + 1^2) = 2.2361$$

$$\text{Cosine} = (A \text{ dot } B) / (||A|| ||B||) = 4 / (2.4495 * 2.2361) = 0.7302383227$$



# Evaluation of Cosine

- Output will receive a value between -1 to 1
- -1: Completely Dissimilar
  - The cat is sleeping
  - The cat is not sleeping
- 1: Completely Similar
  - The cat is sleeping.
  - The kitten is napping.
- 0: Indicates that they are orthogonal (not similar at all)
  - The cat is sleeping
  - The dog is barking





# Guide Questions

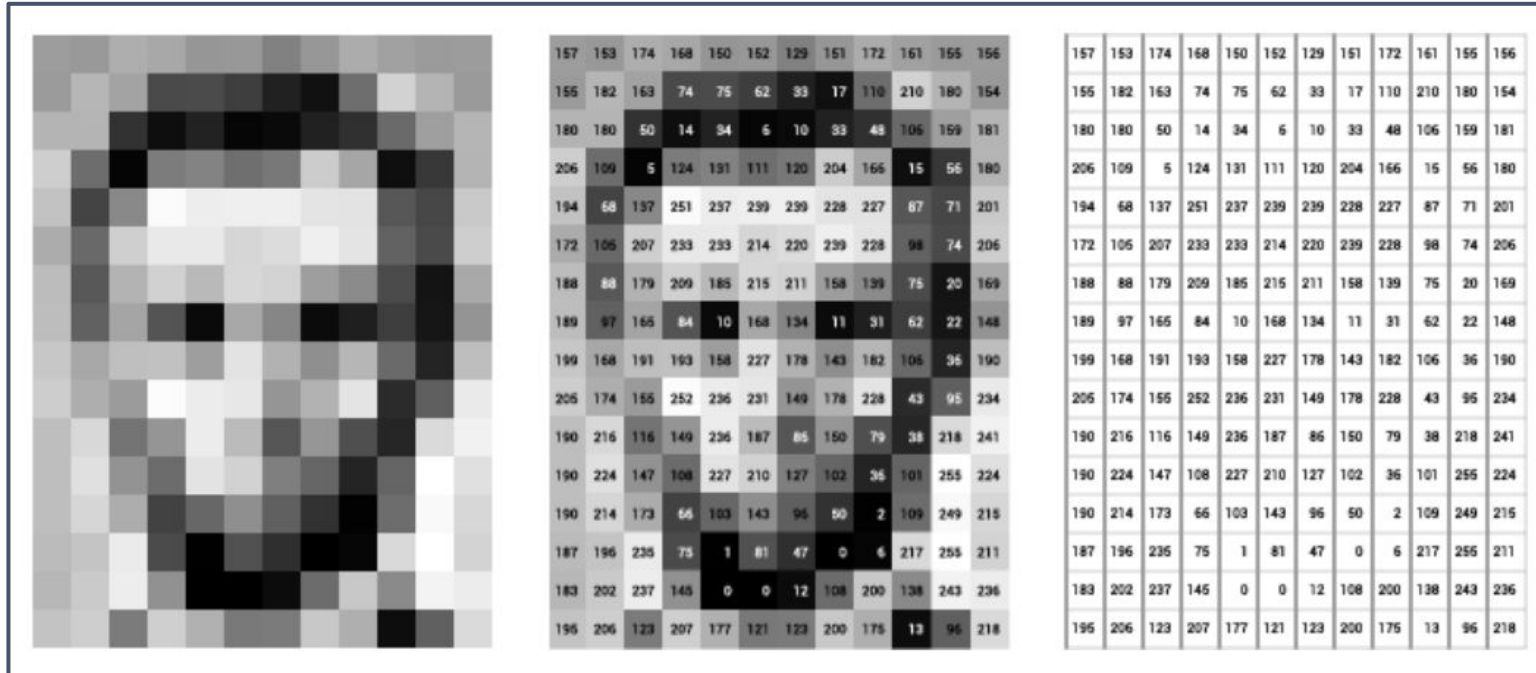
- What are potential problems of cosine similarity based on count?
- Can similarity measures be applied to a probabilistic framework when dealing with text?



# Images



- Pixel Data: Dimensionality = Length x Width





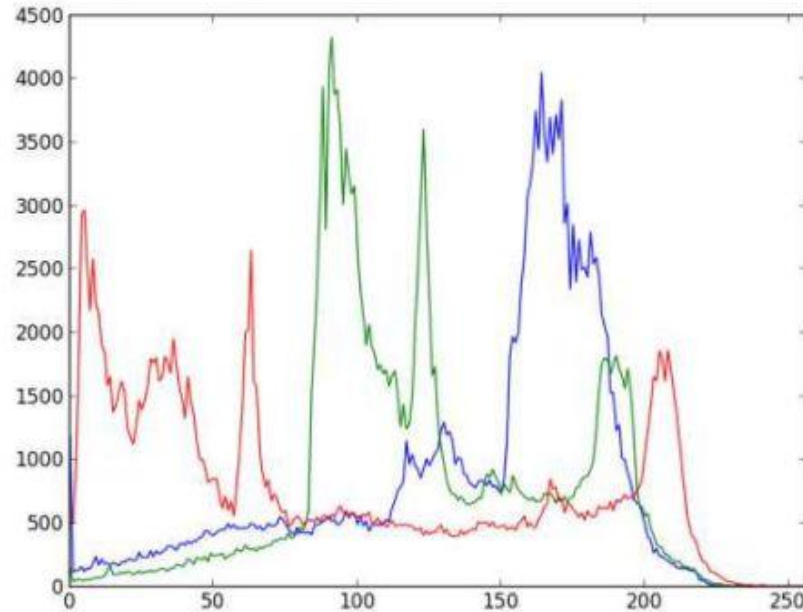
# Global Features

- Characteristics or patterns in an image that describe the entire image as a whole
- Although global features can be represented as multi-dimensional data, it only captures the gist of an image as opposed to local characteristics that compose an image
- Examples of Global Features:
  - Histograms
  - Texture



# Histograms

- Count of similar pixel values
- Question: What are problems with granularity with histograms?





- Gray Level Co-occurrence Matrix

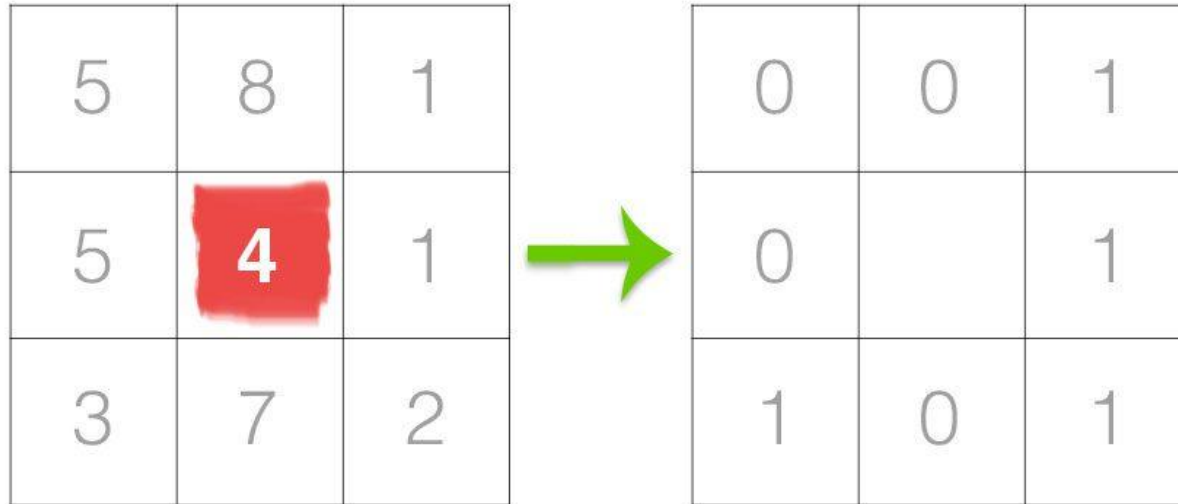
## Statistics Group:

GLCM Mean	$\mu_i = \sum_{i,j=1}^N i(P_{i,j})$ $\mu_j = \sum_{i,j=1}^N j(P_{i,j})$ $\mu = \frac{(\mu_i + \mu_j)}{2}$
GLCM Variance	$\sigma_i^2 = \sum_{i,j=1}^N P_{i,j}(i - \mu_i)^2$ $\sigma_j^2 = \sum_{i,j=1}^N P_{i,j}(j - \mu_j)^2$ $\sigma^2 = \frac{(\sigma_i^2 + \sigma_j^2)}{2}$
GLCM Correlation	$\sum_{i,j=1}^N P_{i,j} \left[ \frac{(i - \mu_i)(j - \mu_j)}{\sqrt{\sigma_i^2 \sigma_j^2}} \right]$



# Local Binary Patterns

- Like GLCM, LBP is used to determine texture features in a given image
- Captures local texture variations by encoding the relationship between a center pixel and its  $n$ th surrounding neighbors
- LBP1:





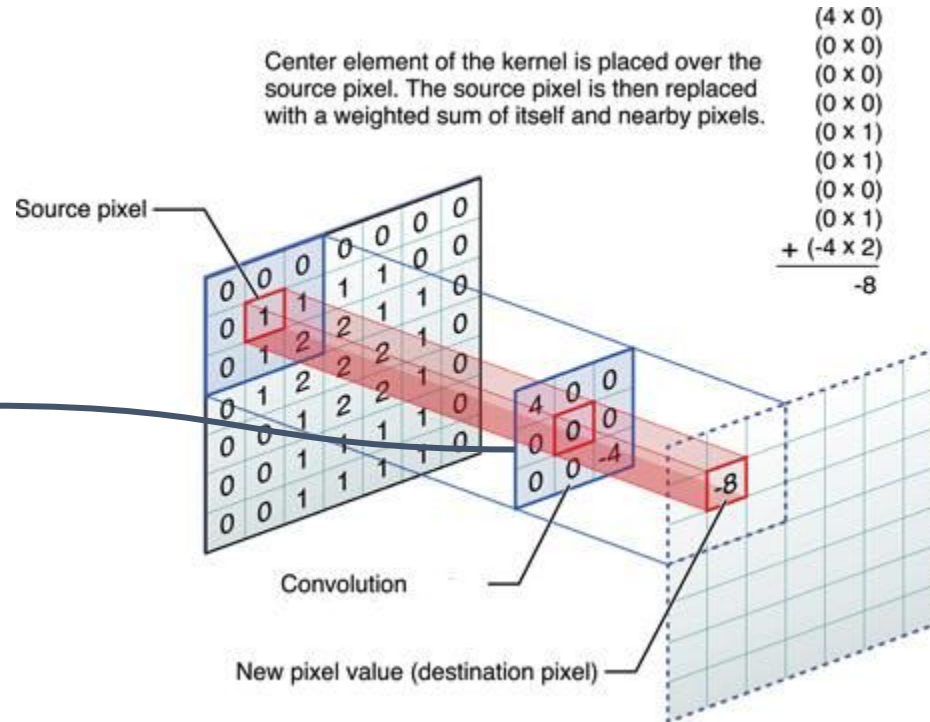
# Deep Learning Features

- High level representation of image data that are learned by deep learning networks
- Creates meaningful patterns, structures and semantic information present in the data
- Semantically Meaningful Features: Encodes information about the content of an image such as the presence (or non-presence) of an object regardless of its location or orientation





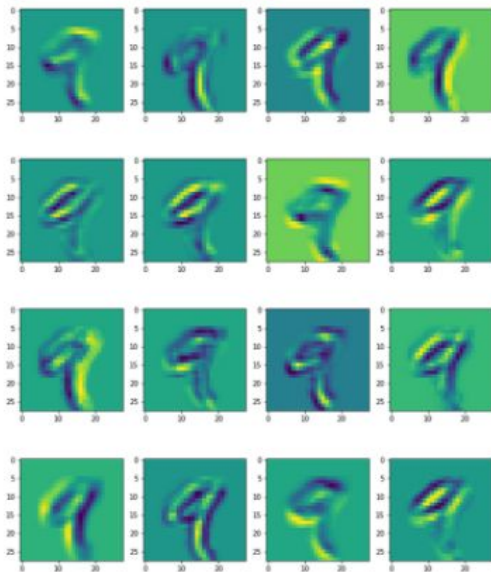
- Convolution Operation



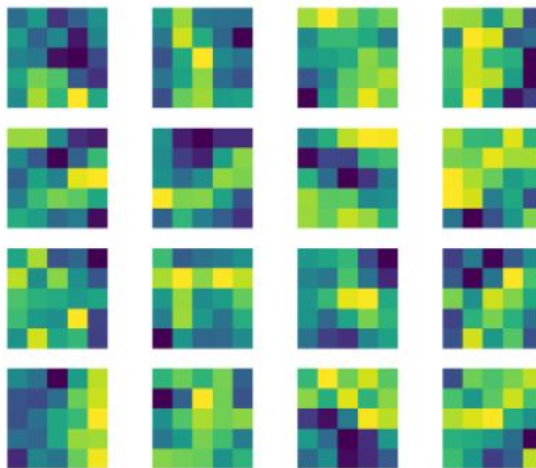


# Convolutional Features

Feature map

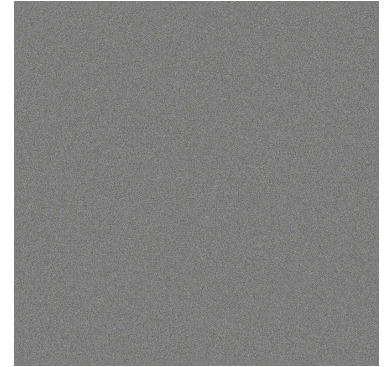
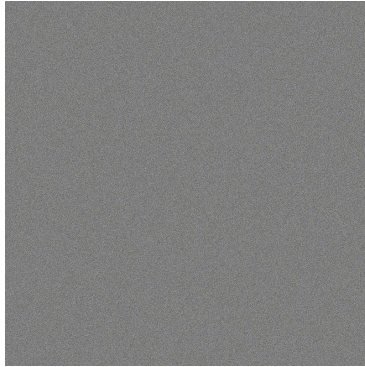


Filters





# Pixel Dreams





# Guide Questions

- What are the potential weaknesses of non-deep learning features?
- What about videos? How do you think videos can be represented as data?



# Audio



# Audio Signal

- Representation of sound waves that can be processed and analyzed by a computer
- A sequence of numerical values that represent the instantaneous amplitude of a sound wave at discrete time points



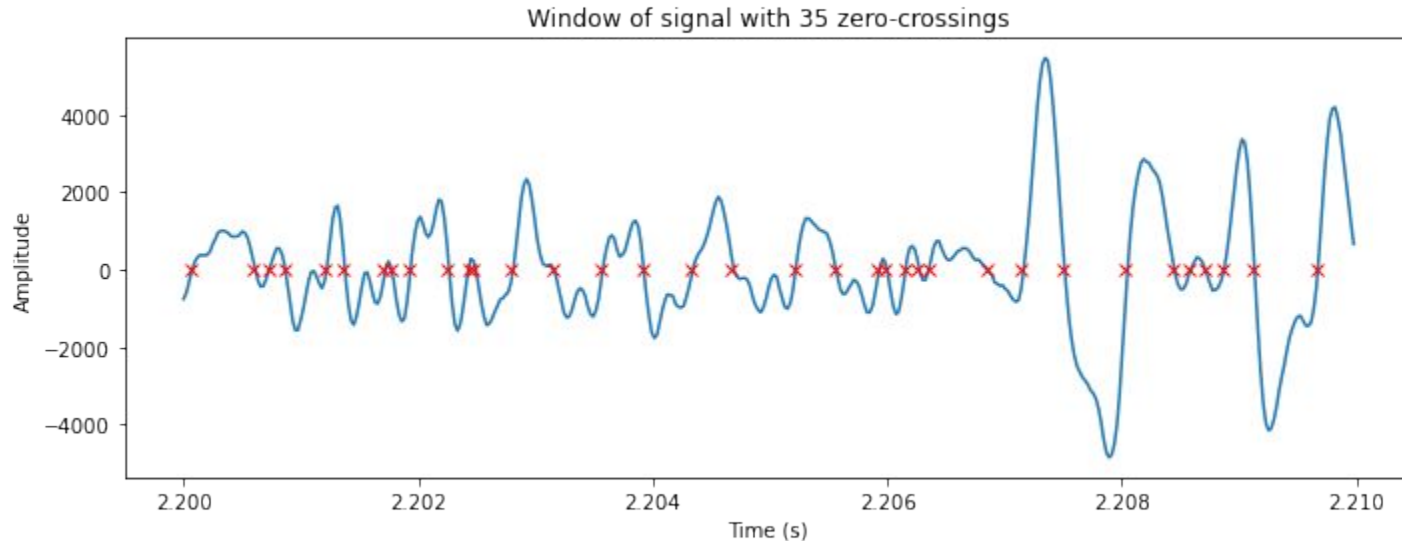
# Temporal Features

- Temporal: A feature in relation to time or some chronological order
- Basic Examples:
  - Zero Cross Rate
  - Root Mean Square Energy
  - Temporal Centroid



# Zero Cross Rate

- Number of times a signal crosses the zero level within a specified time frame
- Characterizes the frequency content and noisiness of a signal







# Root Mean Square Energy

- Energy of an audio signal over time

$$x_{\text{rms}} = \sqrt{\frac{1}{n} (x_1^2 + x_2^2 + \cdots + x_N^2)}$$



# Temporal Centroid

- Provides insight into where the “center” of the audio’s energy distribution is over time.
- Measured by the temporal location of the average energy of an audio signal

```
total_energy = sum(x ** 2 for x in audio_signal)
temporal centroid = sum((i / len(audio signal)) * (x ** 2)
for i, x in enumerate(audio_signal)) / total_energy
```



# Questions?

## Thank you!